

AdSA-ALMD user guide: from ABS Census data to a labour market delineation and back

This is an end-to-end walkthrough for someone who has ABS Census microdata (PLIDA / DataLab) and wants to:

1. turn it into the `w / adj` inputs the `adsalmd` package needs,
2. run the delineation pipeline (`run_adsa_pipeline()`), and
3. convert the result back into an ABS-style correspondence table — the same row-per-region, code/name/ratio layout as an official ABS geography correspondence file — so the output can be read, joined, and reported on the way ABS geography products normally are.

Everything in this guide, and every function it calls from `dev/`, is internal tooling — not part of the installable `adsalmd` package. The package itself (`R/`) only knows about `w`, `adj`, and numeric thresholds; it has no ABS-specific code. `dev/` bridges the gap for this specific data source. See `dev/README.md` for the shorter reference version of steps 1–2 below.

0. Prerequisites

R packages - Base R covers everything except the adjacency step. - `sf` and `spdep`, only for building `adj` from a shapefile.

```
install.packages(c("sf", "spdep"))
```

Data you need in hand

#	What	Where it comes from
1	Person-level Census microdata	PLIDA / DataLab extract, or a CSV with the same columns
2	DZN SA2 concordance	ABS-published <code>DZN_SA2_<year>_AUST.csv</code>
3	SA2 boundary shapefile	ABS ASGS digital boundary files (<code>SA2_<year>_AUST_GDA<datum>.sh</code>)

#	What	Where it comes from
		p)

Source the helper functions (adjust the path to wherever your checkout of this repo lives):

```
source("dev/abs_census_to_od.R")
source("dev/abs_adjacency_from_shapefile.R")
source("dev/validate_lma_inputs.R")
source("dev/lma_result_to_abs_correspondence.R")

library(adsalmd) # the installed package, for run_adsa_pipeline() etc.
```

1. Prepare your Census extract

If your data arrives as plain CSV rather than a DataLab SAS extract, it still needs these exact ABS Census Dictionary column names (or you pass the matching names as arguments — see step-by-step below):

Header	Meaning	Notes
ABSPID	Person identifier	carried through, not filtered on
AGEP	Age	used for the working-age filter
LFSP	Labour force status	1 / 2 / 3 = employed full-time / part-time / away from work
PURP	Place of Usual Residence	already SA2-coded
POWP	Place of Work	DZN -coded, needs conversion (step 1.2)
IFPOWP	Imputation flag for POWP	0 = observed, 1 = imputed

Don't have an extract to hand yet? `dev/sample_data/` ships a small synthetic one (1,000 fake person records, 10 Brisbane SA2s) with these exact column names and coding scheme, plus a matching DZN SA2 lookup, so you can run this entire section — and check the helper functions actually work — before you've pulled a real DataLab extract:

```
census_raw      <- read.csv("dev/sample_data/census_extract_sample.csv", colClasses =
"character")
dzn_sa2_lookup <- read.csv("dev/sample_data/DZN_SA2_2016_AUST_sample.csv", colClasses =
"character")
```

Everything from here through 1.5 runs unchanged on this sample data; see 1.6 below for the full sample walkthrough with expected output. When you're ready to switch to your own data, it's a one-line swap — same column names throughout:

```
census_raw <- read.csv("my_census_extract.csv", colClasses = "character")
```

Read `AGEP` as numeric if it came in as character:

```
census_raw$AGEP <- as.integer(census_raw$AGEP)
```

1.1 Filter to employed, working-age, non-imputed records

```
census_filtered <- filter_census_employed(
  census_raw,
  age_col = "AGEP", age_min = 15, age_max = 64,
  lfsp_col = "LFSP", lfsp_employed = c("1", "2", "3"),
  ifpowp_col = "IFPOWP", require_non_imputed = TRUE
)
```

If your extract already excludes imputed records, or uses different age bounds, adjust the arguments — nothing here is hardcoded.

1.2 Resolve place-of-work from DZN to SA2

```
dzn_sa2_lookup <- read.csv("DZN_SA2_2016_AUST.csv", colClasses = "character")
# expects columns DZN_CODE_2016, SA2_MAINCODE_2016 by default -- rename or
# pass dzn_lookup_col / sa2_lookup_col if your file uses different headers

census_sa2 <- convert_powp_to_sa2(census_filtered, dzn_sa2_lookup)
```

(If you loaded `dzn_sa2_lookup` from `dev/sample_data/` above, skip the `read.csv()` line here — it's already in your environment — and go straight to `convert_powp_to_sa2()`.)

`census_sa2` now has `origin_SA2` (= `PURP`) and `destination_SA2` (from `POWP` via the DZN lookup). Rows whose `POWP` didn't match any DZN in the lookup are dropped by default (`drop_unmatched = TRUE`).

1.3 Fix the canonical SA2 order for this run

`W` and `adj` must agree, row for row, on which SA2 is unit `i`. Decide that order once, here, and reuse it everywhere below:

```
sa2_codes <- sort(unique(c(census_sa2$origin_SA2, census_sa2$destination_SA2)))
```

(Or supply your own study-area list of SA2 codes if you want to include SA2s with zero recorded flow, or exclude some.)

1.4 Build the OD matrix `W`

```
W <- build_od_matrix(census_sa2, sa2_codes)
# W[i, j] = number of employed residents of sa2_codes[i] working in sa2_codes[j]
```

1.5 (Optional but recommended) look at self-containment

This is the same self-containment figure the pipeline's `sc_min` threshold is checked against, so looking at it now tells you what thresholds are even achievable for your data before you spend a run finding out:

```
sc_table <- compute_self_containment(W)
summary(sc_table$SC)
```

1.6 Sample-data walkthrough (expected output)

Running 1–1.5 above against `dev/sample_data/` end to end:

```

source("dev/abs_census_to_od.R")

census_raw      <- read.csv("dev/sample_data/census_extract_sample.csv", colClasses =
"character")
census_raw$AGEP <- as.integer(census_raw$AGEP)
dzn_sa2_lookup <- read.csv("dev/sample_data/DZN_SA2_2016_AUST_sample.csv", colClasses =
"character")

census_filtered <- filter_census_employed(census_raw)
census_sa2      <- convert_powp_to_sa2(census_filtered, dzn_sa2_lookup)
sa2_codes       <- sort(unique(c(census_sa2$origin_SA2, census_sa2$destination_SA2)))

W              <- build_od_matrix(census_sa2, sa2_codes)
sc_table <- compute_self_containment(W)

```

gives you:

```

> nrow(census_raw)
[1] 1000
> nrow(census_filtered)      # after the employed / working-age / non-imputed filter
[1] 400
> dim(W)
[1] 10 10
> sum(W)                     # equals nrow(census_sa2) -- every filtered record landed
somewhere
[1] 400
> summary(sc_table$SC)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.02381 0.03237 0.05583 0.07002 0.11076 0.12963

```

Self-containment is low across the board here because the sample data assigns `PURP` / `POWP` uniformly at random — real commuting data is far more spatially clustered. Don't read anything into the numbers themselves; the point is that the pipeline runs and `W` comes out the right shape. `dev/sample_data/generate_sample_data.R` regenerates both CSVs if you want to tweak the scenario (more SA2s, a different sample size, etc).

2. Build the adjacency matrix

Using an ABS SA2 boundary shapefile and queen contiguity (shared edge or vertex — the ABS/LMA-delineation convention):

```
adj <- build_adjacency_from_shapefile(  
  shapefile_path = "SA2_2016_AUST_GDA2020.shp",  
  sa2_codes       = sa2_codes,      # same vector as step 1.3, same order  
  sa2_id_col      = "SA2_MAIN16",  # check names(sf::st_read(...)) for your ASGS edition  
  queen          = TRUE  
)
```

This reorders the shapefile to `sa2_codes` 's order for you and errors if any SA2 in `sa2_codes` is missing from the shapefile — fix that before continuing rather than silently dropping units.

3. Validate before you run anything expensive

```
qa <- validate_W_adj(W, adj, census_sa2 = census_sa2)  
qa$report          # full check-by-check table  
stopifnot(qa$pass)  # stop here if anything failed
```

Read every `FAIL` line before proceeding — in particular, “W and adj have identical SA2 order” failing means `run_adsa_pipeline()` will silently delineate the wrong geography (right matrices, mismatched labels).

4. Run the pipeline

```

row_W <- rowSums(W)
col_W <- colSums(W)

res <- run_adsa_pipeline(
  W = W, adj = adj, row_W = row_W, col_W = col_W,
  SC_min = 0.55,      # from sc_table above -- pick a threshold most SA2s can plausibly
meet
  Pop_min = 5000,     # minimum viable labour market population
  Pop_max = 50000,    # maximum, or leave as Inf
  r = 3,
  L = 20, l = 10, T0_samples = 20, # small first: size the run before scaling up
  seed = 2026
)

```

Start small (`L = 20` or so) to see how long a trial takes on your data before committing to the default `L = 1000`. For a final delineation, don't rely on one seed — run several and compare (see `run_adsa_pipeline_parallel()`; the spread across seeds tells you whether the search found one stable structure or several competing ones).

What you get back

- `res$init_sol` — SARA's feasible starting partition (before annealing).
- `res$best_sol` — the final partition **after** annealing and misallocation refinement. Same length and order as `sa2_codes`: `best_sol[i]` is the LMA label of `sa2_codes[i]`.
- `res$lma_metrics` — one row per LMA: `cluster`, `pop`, `scss`, `scds`, `sc`, `n_bgus`.
- `res$gci`, `res$base_fitness`, `res$n_clusters` — recomputed for `best_sol`.
- `res$best_fitness`, `res$history`, `res$elite`, `res$op_tracker` — describe the *pre-refinement* search (see `?run_adsa_pipeline` for why these don't quite match `best_sol`).

Re-check feasibility — stage 3 (refinement) doesn't enforce population bounds:

```
with(res$lma_metrics, all(pop >= 5000 & pop <= 50000))
```

5. Convert the result back into ABS correspondence format

ABS geography correspondence files (e.g. “SA2 (2016) to SA4 (2016) Correspondence”) are one row per lower-level region, with the higher-level region's code and name, and `RATIO_FROM_TO` / `RATIO_`

TO_FROM columns. `build_sa2_lma_correspondence()` produces the same layout for your delineation, treating each LMA as the “higher-level region”:

```
sa2_names <- setNames(my_sa2_lookup$SA2_NAME_2016, my_sa2_lookup$SA2_MAINCODE_2016)

corr <- build_sa2_lma_correspondence(
  sol          = res$best_sol,
  sa2_codes    = sa2_codes,
  sa2_names    = sa2_names[sa2_codes], # optional; omit if you don't have names
  row_W        = row_W,                # population proxy for RATIO_TO_FROM
  lma_code_prefix = "LMA",
  sa2_year     = "2016", lma_year = "2016"
)

head(corr)
#   SA2_MAINCODE_2016 SA2_NAME_2016 LMA_CODE_2016          LMA_NAME_2016 RATIO_FROM_TO
RATIO_TO_FROM
# 1           301011001 Fortitude Valley          LMA001 Local Labour Market Area 1
1           0.34
# ...
```

- `** RATIO_FROM_TO **` is always `1` here: this is a strict partition, so 100% of each SA2 belongs to exactly one LMA (unlike some ABS correspondences between geographies that don't nest cleanly).
- `** RATIO_TO_FROM **` is this SA2's share of its LMA's total population (or of `row_W`, if you didn't supply an actual population vector) — pass a real ERP population vector via `population =` for a more accurate figure than the commuting-flow proxy.

If you supplied real SA2 population instead of the `row_W` proxy:

```
corr <- build_sa2_lma_correspondence(
  sol = res$best_sol, sa2_codes = sa2_codes, sa2_names = sa2_names[sa2_codes],
  population = my_sa2_population # named or positional, aligned to sa2_codes
)
```

5.1 (Optional) attach per-LMA quality metrics

To carry `pop / sc` etc. from `res$lma_metrics` alongside each SA2's assignment (handy for a single reportable table):

```
corr_full <- attach_lma_metrics(corr, res$best_sol, res$lma_metrics)
```

5.2 Export

```
write.csv(corr_full, "SA2_to_LMA_2016_Correspondence.csv", row.names = FALSE)
```

That CSV is now in the same shape analysts and downstream tools expect from an ABS correspondence product — joinable to any other SA2-keyed dataset, and readable without any `ad` `salmd` -specific knowledge.

6. Full worked example (one script)

```

source("dev/abs_census_to_od.R")
source("dev/abs_adjacency_from_shapefile.R")
source("dev/validate_lma_inputs.R")
source("dev/lma_result_to_abs_correspondence.R")
library(adsalmd)

census_raw      <- read.csv("my_census_extract.csv", colClasses = "character")
census_raw$AGEP <- as.integer(census_raw$AGEP)
dzn_sa2_lookup  <- read.csv("DZN_SA2_2016_AUST.csv", colClasses = "character")

census_filtered <- filter_census_employed(census_raw)
census_sa2      <- convert_powp_to_sa2(census_filtered, dzn_sa2_lookup)
sa2_codes       <- sort(unique(c(census_sa2$origin_SA2, census_sa2$destination_SA2)))

W  <- build_od_matrix(census_sa2, sa2_codes)
adj <- build_adjacency_from_shapefile("SA2_2016_AUST_GDA2020.shp", sa2_codes)

qa <- validate_W_adj(W, adj, census_sa2 = census_sa2)
stopifnot(qa$pass)

row_W <- rowSums(W); col_W <- colSums(W)

res <- run_adsa_pipeline(
  W = W, adj = adj, row_W = row_W, col_W = col_W,
  SC_min = 0.55, Pop_min = 5000, Pop_max = 50000, r = 3,
  L = 1000, seed = 2026
)

corr <- build_sa2_lma_correspondence(res$best_sol, sa2_codes, row_W = row_W)
full <- attach_lma_metrics(corr, res$best_sol, res$lma_metrics)
write.csv(full, "SA2_to_LMA_2016_Correspondence.csv", row.names = FALSE)

```

7. Troubleshooting

Symptom	Likely cause
<code>validate_W_adj()</code> fails "W and adj have identical SA2 order"	You changed or re-sorted <code>sa2_codes</code> between calling <code>build_od_matrix()</code> and <code>build_adjacency_from_shapefile()</code> . Use the exact same vector for both.

Symptom	Likely cause
<code>build_adjacency_from_shapefile()</code> errors "sa2_codes not found in shapefile"	Wrong <code>sa2_id_col</code> for your ASGS edition (check <code>names(sf::st_read(path))</code>), or your <code>sa2_codes</code> includes SA2s outside the shapefile's coverage (e.g. a different state/edition).
<code>convert_powp_to_sa2()</code> drops most rows	Wrong <code>dzn_lookup_col</code> / <code>sa2_lookup_col</code> , or a DZN/SA2 vintage mismatch between the Census extract and the lookup file.
All <code>SC</code> values in <code>compute_self_containment()</code> are low	Your SA2s might be too small relative to real commuting patterns for the chosen study area — expect to need SARA/annealing to aggregate several SA2s per LMA; don't set <code>SC_min</code> above what's achievable pre-aggregation.
<code>lma_metrics</code> shows populations outside <code>Pop_min</code> / <code>Pop_max</code> after <code>run_adsa_pipeline()</code>	Expected — stage 3 (refinement) doesn't enforce population bounds (see <code>?run_adsa_pipeline</code> , "Details"). Re-check and treat out-of-bounds LMAs as needing another pass or manual review.
Correspondence <code>RATIO_TO_FROM</code> looks wrong/uniform	You didn't pass <code>population</code> or <code>row_w</code> , so every SA2 was weighted equally within its LMA (a warning is printed when this happens).

8. Where each piece lives

Step	Function	File
Filter microdata	<code>filter_census_employed()</code>	<code>dev/abs_census_to_od.R</code>
DZN SA2	<code>convert_powp_to_sa2()</code>	<code>dev/abs_census_to_od.R</code>
Build <code>w</code>	<code>build_od_matrix()</code>	<code>dev/abs_census_to_od.R</code>
Self-containment diagnostics	<code>compute_self_containment()</code>	<code>dev/abs_census_to_od.R</code>
Build <code>adj</code>	<code>build_adjacency_from_shapefile()</code>	<code>dev/abs_adjacency_from_shapefile.R</code>
QA <code>w</code> / <code>adj</code>	<code>validate_w_adj()</code>	<code>dev/validate_lma_inputs.R</code>

Step	Function	File
Run the pipeline	<code>run_adsa_pipeline()</code> / <code>run_adsa_pipeline_parallel()</code>	public package (<code>R/11_pipeline.R</code>)
Result correspondence table	<code>build_sa2_lma_correspondence()</code>	<code>dev/lma_result_to_abs_correspondence.R</code>
Attach LMA metrics	<code>attach_lma_metrics()</code>	<code>dev/lma_result_to_abs_correspondence.R</code>